

Capture your best planner's knowledge

before it walks out the door.

An advisory AI agent for production scheduling and JIT procurement — runs on a single machine inside your plant, learns from every decision your planner makes, and turns 15 years of tribal knowledge into an asset you own.



Galvon · 2026

The Problem

Every manufacturing company has **one or two people who really know how to run the place.**

They decide what to make today. What to order. From whom. When to expedite. When to wait. Their knowledge is in their head — built from 15 years of being burned by Supplier A in monsoon, knowing Machine 4 hates a 2-hour run-then-stop, knowing Tata pays a ₹50K/day penalty if you slip.

When that person is on leave, takes a competing job, or retires, the company loses money.

What This Costs You — Today

A typical mid-size plant loses ₹2–8 Cr a year to scheduling mistakes nobody captures.

WHERE THE MONEY GOES	TYPICAL ANNUAL LOSS
Late-delivery penalties + lost orders	₹40 L - 2 Cr
Expedited freight to recover slippage	₹20 - 80 L
Scrap from bad changeover sequencing	₹30 L - 1.5 Cr
Lost shifts to setup/changeover thrash	₹50 L - 2 Cr
Excess working capital in safety stock	₹1 - 3 Cr tied up

These numbers don't show up in any one report. They're hidden across penalty debits, freight invoices, scrap registers, and inventory holding cost — and they grow when your best planner is on leave.

The Agent in One Picture



The agent does **not** decide for you. It recommends. The planner approves. Every edit teaches it. Everything is auditable.

Two Vaults — Playbook and Logbook

The agent reads from a **Skills Vault** (the factory's playbook, written by humans) and writes to a **Memory Vault** (the logbook of every decision, correction, and lesson learned).

Skills Vault — your playbook


- Written by you and the planner
- Plain markdown files in git
- VIP customer rules, machine quirks, supplier policies, changeover rules
- Read-only to the agent — it cannot accidentally edit your rules

Memory Vault — the agent's logbook

- Written by the agent every time a decision happens
- Every recommendation, every edit, every reason
- After repeated overrides, the agent drafts a **Lesson** for you to approve
- 100% yours, in plain markdown, in git

Both are plain text. Both are yours. If we disappear tomorrow, your knowledge is still readable in Notepad.

How It Learns — The Correction Loop



This is the IP. Every override is a permanent training signal. Three similar overrides become a promoted rule. The agent grows without anyone writing code.

A Tuesday Morning at "Acme Auto Parts"

9:00 AM — Rakesh arrives. The agent has already drafted today's schedule and tomorrow's reorder list, with citations to every rule it used.

9:05 AM — Rakesh reviews. The Tata priority rule is right. The Machine 4 warmup rule is right. But he bumps a Mahindra work order ahead because Mahindra called him last night about a launch event Friday. He types one sentence — "*Mahindra emergency, Friday launch*" — and clicks Approve.

9:06 AM — The agent writes a correction note: what it said, what Rakesh did, why.


11:30 AM — Machine 4 trips. The agent re-plans in seconds and offers two options with citations. Rakesh picks one. No correction needed this time.

30 DAYS LATER — Rakesh has overridden the agent 4 times on Mahindra emergencies. The agent drafts a Lesson:

GALVIN **MAN MATHURING AT**
"When a Mahindra emergency call is logged, treat Mahindra orders as same-
priority as Tata for 48 hours."

The Value Compounds

Planner
override rate



PHASE	WHAT'S HAPPENING
Week 1	Agent ships generic. ~80% override rate. Corrections are the product, not the recommendations.
Month 1	~50 corrections captured. Override rate ~50%.
Month 3	~10 promoted Lessons. Override rate ~25%.
Month 6	~30 active Lessons. Override rate ~10%. The agent recommends what your best planner would, even when they're on leave

What Makes Real Manufacturing Hard

Most "AI scheduling" demos work on toy plants — one product, one line, one supplier. **Real plants don't look like that.** Your products share upstream processes, your BOMs are multi-level, and one quality reject cascades into a dozen downstream orders.

REAL-PLANT PROBLEM	WHAT THE AGENT DOES
Multiple finals share an upstream WIP (e.g. one furnace heat feeds rebar + wire rod + structural)	Models items as a DAG, not a flat list. Aggregates demand at the shared step.
Multi-level BOM with shared sub-assemblies	BOM-explode walks the tree, computes cumulative lead time, hits MOQ once instead of three small batches.
Supply at a shared step < total demand from finals	Runs your allocation policy as a skill (which customer takes the hit, which slips, why).
Quality reject at a shared upstream step	Walks where_used forward, shows you which downstream orders will slip and by how much.

One Engine, Three Verticals

Same core. Three skill packs. The agent ships pre-loaded with rules that match how your industry actually runs.

Discrete

Auto parts · electronics · panel building · machining

- Kanban-pull triggers and WIP caps
- Due-date prioritization with critical-ratio tweaks
- Lot-traceability protection on serialized parts
- Cell-grouping and changeover sequencing

Batch

Chemicals · food · pharma · paints · cosmetics

- Recipe-driven material consumption
- Campaign batching, grade transitions, CIP/SIP rules
- Expiry-aware FIFO
- 21-CFR-11 / GMP-aware scheduling

Process

Cement · steel · paper · refineries · smelters

- TOC drum-buffer-rope bottleneck protection
- Grade-transition slope rules
- Utility-tariff peak/off-peak load shifting
- Variable-yield JIT for consumables

Your Data Stays In Your Plant

WHERE IT RUNS	WHAT IT DOES	COST
Local Ollama (on your plant PC)	90% of decisions — routine scheduling, JIT triggers, retrieval, Q&A	Hardware amortization only
Claude API (escalation only)	Complex re-planning, lesson synthesis, anomaly explanation	~₹250-400 / month per plant
Haiku classifier	Routes each query to the right model	~₹40 / month per plant

Minimum hardware to start: 16 GB RAM, 8 cores. No GPU required for Bronze tier.

The agent runs **locally**. Routine queries — *"recompute the kanban refill"*, *"explode this BOM"* — never leave your network. Only the genuinely hard reasoning escalates to a remote model, and even then no master data crosses the wire.

The plant never stops getting useful output, even when the internet is down. Hard queries queue with a clear status. Routine schedules and JIT triggers stay current.

Deployment Roadmap — Bronze → Platinum

You don't commit to everything on Day 1. We tier the engagement so you see value in Week 1 and earn the right to escalate.

TIER	WHAT YOU GIVE US	WHAT THE AGENT DOES	TIME TO VALUE
Bronze	Item master + BOM + supplier list + weekly Excel exports	Weekly schedule recommendation + JIT reorder list + generic priority ranking	Week 1
Silver	Above + ERP read or daily file exports + 15–20 customer skills authored in a 1-day workshop	Daily schedule, daily PO suggestions, supplier-reliability-aware JIT, customer-aware prioritization	Week 4
Gold	Above + MES/SCADA live feed + production confirmations + GRN data	Reactive rescheduling on machine-down, closed-loop learning, promoted Lessons, anomaly explanations	Month 3
Platinum	Time-gated, not feature-gated. 6+ months of accumulated corrections and Lessons.	Genuinely personalized. Override rate ~10%. Defensible knowledge moat that competitors cannot port.	Month 6+

What You Give Us / What You Own

Minimal intake — Week 1

- Item master + BOM + supplier list (Excel is fine)
- One named planner — 30 min/day for 60 days
- One IT person — 4 hours total, Week 1
- One shared folder for daily file drops
- One PC inside the plant: 16 GB RAM, 8 cores

That's it. No ERP rip-and-replace. No MES integration on Day 1. No 6-month "data foundation" prerequisite.

What you keep — forever

- **Skills Vault** — every rule you authored
- **Memory Vault** — every decision, correction, Lesson the agent learned
- Both are plain markdown files in your own git repo
- If you cancel the subscription, the vaults stay yours
- A successor planner starts at Month-6 quality, not Week-1

The agent runtime is the recurring revenue. The vaults are the customer-owned asset. That asymmetry is the deal.

The Next Step

We don't sell on the first call. We run a **30-minute Customer Readiness Audit** — a checklist against your master data, your planner availability, and your IT access.

You walk away with:

1. A **tier recommendation** (Bronze / Silver / Gold) based on what your plant can realistically support today.
2. A **go / no-go** call, given honestly. We will tell you if you should fix your data foundation before starting.
3. A **fixed-fee Week-1 pilot quote**, if you proceed.

***Book the audit:** chatgptnotes@gmail.com · Dr. Murali B. K. — Galvon Or talk to Saikat (automation engineer, 20+ years) directly.*